

Project 1: numerically solving a differential equation through a linear system

Daniele Brugnara <daniebru@mail.uio.no>
Martina Cammilli <martinacam@mail.uio.no>
Matteo Seclì <mattes@mail.uio.no>

September 8, 2014

Abstract

In this project we aim to solve a special kind of differential equation using a numerical procedure that allows us to express the equation through a linear system. We will study some algorithms to solve such a problem, focusing on the efficiency of the program, setting our goal more on speed than generality.

1 Introduction

The differential equation we're interested in studying is of the type

$$u''(x) = -f(x) \quad (1)$$

In our case we will limit our solutions using the contour conditions of $u(0) = 0$ and $u(L) = 0$, where $[0, L]$ is our domain of integration. Using Taylor expansion it is possible to express the second derivative of a function $u(x)$ as

$$u''(x) = \frac{u(x-h) - 2u(x) + u(x+h)}{h^2} + \mathcal{O}(h^2) \quad (2)$$

It is possible to discretize equation (1) using N points, obtaining:

$$u_i'' = \frac{u_{i-1} - 2u_i + u_{i+1}}{h^2} = -f_i \quad i \in \{1 \dots N\}$$

Using the matrix representation, we can write equation (1) as a system of linear equations

$$\begin{pmatrix} 2 & -1 & 0 & 0 & \dots & 0 \\ -1 & 2 & -1 & 0 & \dots & 0 \\ 0 & -1 & 2 & -1 & \dots & 0 \\ \vdots & \vdots & & \ddots & & \vdots \\ 0 & 0 & \dots & -1 & 2 & -1 \\ 0 & 0 & \dots & \dots & -1 & 2 \end{pmatrix} \begin{pmatrix} u_0 \\ u_1 \\ u_2 \\ \vdots \\ u_{N-2} \\ u_{N-1} \end{pmatrix} = h^2 \begin{pmatrix} f_0 \\ f_1 \\ f_2 \\ \vdots \\ f_{N-2} \\ f_{N-1} \end{pmatrix} \quad (3)$$

Note how, with this system it is already implied that $f(0) = 0$ e $f(L) = 0$, since the first and last equations state that

$$\begin{aligned} h^2 f_0 &= 2u_0 - u_1 = -1u_{-1} + 2u_0 - u_1 \\ h^2 f_{N-1} &= -u_{N-2} + 2u_{N-1} = -u_{N-2} + 2u_{N-1} - u_N \end{aligned}$$

This follows from the boundary conditions of the differential equations that state that $u_{-1} = u(0) = 0$ and $u_N = u(L) = 0$.

This linear system is indeed very particular and has a clear pattern. We will first focus on finding a solving algorithm for a general tridiagonal matrix and after we will try to implement another program to solve this particular system with the intent of lowering the number of calculation and therefore the computation time.

2 General algorithm for solving a tridiagonal matrix through back and forward substitution

A general tridiagonal system can be expressed as

$$\begin{pmatrix} b_0 & c_0 & 0 & 0 & \cdots & 0 \\ a_1 & b_1 & c_1 & 0 & \cdots & 0 \\ 0 & a_2 & b_2 & c_2 & \cdots & 0 \\ \vdots & \vdots & & \ddots & & \vdots \\ 0 & 0 & \cdots & a_{N-2} & b_{N-2} & c_{N-2} \\ 0 & 0 & \cdots & \cdots & a_{N-1} & b_{N-1} \end{pmatrix} \begin{pmatrix} u_0 \\ u_1 \\ u_2 \\ \vdots \\ u_{N-2} \\ u_{N-1} \end{pmatrix} = h^2 \begin{pmatrix} f_0 \\ f_1 \\ f_2 \\ \vdots \\ f_{N-2} \\ f_{N-1} \end{pmatrix} \quad (4)$$

We will describe the algorithm we used for this system first for a 3×3 tridiagonal matrix, and after we will demonstrate its validity for a square tridiagonal matrix of whatever dimension (proof by induction). We will also limit to the case $c_0 = c_1 = \dots = c_{N-2}$, $a_0 = a_1 = \dots = a_{N-2}$ and $b_0 = b_1 = \dots = b_{N-1}$.

- 3×3 matrix:

$$\begin{aligned} \left(\begin{array}{ccc|c} b & c & 0 & f_0 \\ a & b & c & f_1 \\ 0 & a & b & f_2 \end{array} \right) &\xrightarrow{\text{Passage 1}} \left(\begin{array}{ccc|c} 1 & c/b & 0 & f_0/b \\ a & b & c & f_1 \\ 0 & a & b & f_2 \end{array} \right) \\ &\xrightarrow{\text{Passage 2}} \left(\begin{array}{ccc|c} 1 & c/b & 0 & f_0/b \\ 0 & \frac{b-(c/b)a}{b-(c/b)a} & \frac{c}{b-(c/b)a} & \frac{f_1-af_0}{b-(c/b)a} \\ 0 & 0 & \frac{b-\frac{ac}{b-(c/b)a}}{b-(c/b)a} & \frac{f_2-af_1}{b-(c/b)a} \end{array} \right) = \left(\begin{array}{ccc|c} 1 & c/b & 0 & f_0/b \\ 0 & 1 & \frac{c}{b-ac/b} & \frac{f_1-af_0}{b-(c/b)a} \\ 0 & 0 & 1 & \frac{f_2-af_1}{b-\frac{ac}{b-(c/b)a}} \end{array} \right) \\ &\xrightarrow{\text{Passage 3}} \left(\begin{array}{ccc|c} 1 & 0 & 0 & f_0/b - f_1(c/b) \\ 0 & 1 & 0 & \frac{f_1-af_0}{b-(c/b)a} - f_2 \frac{c}{b-ac/b} \\ 0 & 0 & 1 & \frac{f_2-af_1}{b-\frac{ac}{b-(c/b)a}} \end{array} \right) \end{aligned}$$

Now it's very simple to solve the system.

- $(n+1) \times (n+1)$ matrix.

Before starting to demonstrate that the above passages can be done also for a $(n+1) \times (n+1)$ matrix, supposed that they work for a $n \times n$ one, we can notice that, in general, for a square matrix of optional dimension n , doing the passage 2 until the penultimate row we obtain:

$$A_{n-2,n-1} = \frac{c}{bet(n-2)}$$

(we are using the convention according to with the indexes start from 0, so the last row and column are labelled $n-1$) where

$$bet(m) = b_0 - \frac{ac}{b_1 - \frac{ac}{b_2 - \frac{ac}{\ddots b_{m-1} - \frac{ac}{b_m}}}}$$

(here all the b_i 's have the same value; the index i helps only to count them).

Now we do the passage 2 until the last row (we focus only on the tridiagonal matrix; if we manage to obtain the unitary matrix the system is solved); we obtain:

$$\begin{pmatrix} 1 & c/b & 0 & 0 & \cdots \\ 0 & 1 & \frac{c}{b-ac/b} & 0 & \cdots \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ \cdots & \cdots & \cdots & 1 & c/bet(n-1) \\ \cdots & \cdots & \cdots & 0 & 1 \end{pmatrix} \quad (5)$$

and simply subtracting, from the $(n-1)$ -row, the n -row multiplied for $c/bet(n-1)$:

$$\begin{pmatrix} 1 & c/b & 0 & 0 & \cdots \\ 0 & 1 & \frac{c}{b-ac/b} & 0 & \cdots \\ \cdots & \cdots & \cdots & \cdots & \cdots \\ \cdots & \cdots & \cdots & 1 & 0 \\ \cdots & \cdots & \cdots & 0 & 1 \end{pmatrix} \quad (6)$$

Now, ignoring the n -row and the n -column, we have a $n \times n$ matrix (note again that the indexes start from 0) which we can bring back to the identity going on with the passage 3.

- Algorithm in C++

Translating the above algorithm in C++ language and working only on the vector f and on the solution vector u , we obtain the following code:

```

    u[0]=f[0]/(bet=b);
    for(int j = 1; j < N; j++) {
        gam[j]=c/bet;
        bet=b-a*gam[j];
        u[j]=(f[j]-a*u[j-1])/bet;
    }
    for (int j = (N-2); j >= 0; j--) u[j] -= gam[j+1]*u[j+1];

```

2.1 Particular algorithm

Using the regular Gaussian elimination algorithm we proceed to find a specific solution of our system as follows:

$$\left(\begin{array}{cccccc|c} 2 & -1 & 0 & 0 & \cdots & 0 & f_0 \\ -1 & 2 & -1 & 0 & \cdots & 0 & f_1 \\ 0 & -1 & 2 & -1 & \cdots & 0 & f_2 \\ \vdots & \vdots & & \ddots & & \vdots & \vdots \\ 0 & 0 & \cdots & -1 & 2 & -1 & f_{N-2} \\ 0 & 0 & \cdots & \cdots & -1 & 2 & f_{N-1} \end{array} \right) \rightarrow \left(\begin{array}{cccccc|c} 2 & -1 & 0 & 0 & \cdots & 0 & f_0 \\ 0 & 3 & -2 & 0 & \cdots & 0 & 2f_1 + f_0 \\ 0 & 0 & 4 & -3 & \cdots & 0 & 3f_2 + 2f_1 + f_0 \\ \vdots & \vdots & & \ddots & & \vdots & \vdots \\ 0 & 0 & \cdots & 0 & N & -(N-1) & (N-1)f_{N-2} + \sum_{j=0}^{N-3} (j+1)f_j \\ 0 & 0 & \cdots & \cdots & 0 & N+1 & \sum_{j=0}^{N-1} (j+1)f_j \end{array} \right) \quad (7)$$

We therefore know ahead of times the explicit form of the matrix in upper triangular form and are able to compute all the constant terms of the system as follows

$$\tilde{f}_i = \sum_{j=0}^i (j+1)f_j$$

However, we don't need to compute the sum every time, since we can compute \tilde{f}_i knowing \tilde{f}_{i-1} :

$$\tilde{f}_i = (i+1)f_i + \tilde{f}_{i-1}$$

Once this forward computations are completed, it is possible to proceed with a back substitution, knowing that

$$u_{N-1} = \frac{1}{N+1} \tilde{f}_{N-1}$$

we are able to find the vector of solutions u_i

$$u_i = \frac{1}{i+2} (\tilde{f}[i] + (i+1)u[i+1])$$

Translating the algorithm in C++ code we obtain the following cycle.

```

for (i=1; i!=n; i++) {
    f[i]=(i+1)*f[i]+f[i-1];
}
u[n-1]=f[n-1]/(n+1);
for (i=n-2; i>=0; i--) {
    u[i]=(f[i]+(i+1)*u[i+1])/(i+2);
}

```

In order to obtain less operations, we have implemented in the following code a shift of index. The difference, however was marginally appreciable and barely noticeable (around 1% to 2% relative increase of speed).

```

1  for(int j = 1; j < N; j++) f[j]=(j+1)*f[j]+f[j-1];
2  u[N-1] = f[N-1]/(N+1);
3  int prev_idx = N;
4  for(int j = N - 1; j > 0; j--) {u[j-1]=(f[j-1]+j*u[j])/prev_idx; prev_idx = j;}

```

Note that in the code provided here the indexes are shifted in order to reach a theoretical number of FLOPS equal to $6N$. Further we will discuss better about FLOPS in our program.

2.2 Other algorithms

We have also computed the solution using other algorithms, namely the LU decomposition and the standard Gaussian elimination. Both of these apply to general matrices and therefore are expected to perform not as well as the other ones. We will observe such differences in the next paragraph, when we will compare the different elapsed times. For this algorithms we took advantage of an external library, *armadillo*, as can be seen in the script in Section 7.

3 Solution of the differential equation

Given our differential equation

$$u''(x) = -f(x)$$

it is easy to solve by just integrating the function $f(x)$ and using the boundary condition to determine the value of the two integration constants. Our function is $f(x) = 100e^{-10x}$

$$u'(x) = \int -100e^{-10x} dx = 10e^{-10x} + C_1 \quad u(x) = \int u'(x) dx = -e^{-10x} + C_1 x + C_2$$

and our boundary conditions are $u(0) = -1 + C_2 = 0$ and $u(1) = -e^{-10} + C_1 + C_2 = 0$. Solving the system we obtain $C_1 = e^{-10} - 1$, $C_2 = 1$, which substituted in our family of solutions yields to our analytical solution

$$u(x) = 1 - (1 - e^{-10})x - e^{-10x}$$

This function will be used to calculate the errors caused by our numerical approximation. In Figure 1 we can see the graph of this function as computed by solving the differential equation. Not much can be said except that the boundary conditions are satisfied.




Figure 1: Graph of the solution of the differential equation, computed with 10^4 points, using the special algorithm.

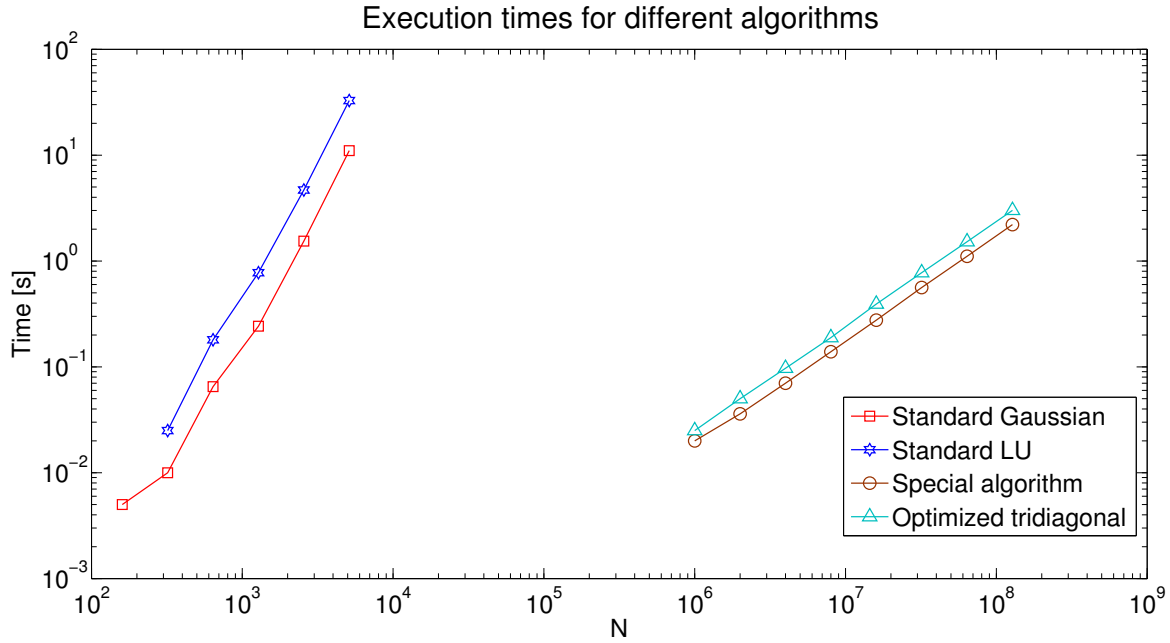


Figure 2: Elapsed time during calculation for the various algorithms on a logarithmic grid.

4 Time

We will now study how the program performs as a function of the dimension of the matrix.

In Figure 2 we can notice how a specialized algorithm is indeed able to significantly cut down the time required for the computation of the solution. In fact, we can easily notice how, with a standard Gaussian elimination, the required time for solving a 2×10^2 dimensional matrix is the same it takes for the specialized algorithm to compute 10^6 points. This yields clearly to a much higher degree of precision and a better deployment of resources.

As a matter of fact, counting the number of operations contained in the specialised algorithm is $6N$, whereas for the LU decomposition it should amount to $N^2 + N^3$. Unfortunately we are unable to appreciate this cubic function, given the scarcity of data, however the slope of the curve is much higher. Moreover the 'optimized tridiagonal' is characterized by $8N$ operations and we are able to notice a clear but not overwhelming difference.

It is also possible to notice how the Gaussian elimination managed to be faster than the Standard LU, this is probably due to the fact that our particular matrix was composed by a really high number of zeroes and armadillo was supposedly able to avoid unnecessary operations among zeroes, thus reducing the computation time.

It is worth compare the "theoretical" FLOPS of the specialised algorithm with the real ones. For $N = 10^8$, the program takes about 1.77 s. All of our tests were conducted on a Ubuntu/Linux machine, equipped with a Intel i7-3610QM processor and 8GB of DDR3 memory. On the page <http://browser.primatelabs.com/geekbench2/2343311> there is a performance table of the processor, and we see that operating with single-core floating point numbers the processor reaches 1.68 Gflops. Considering that we were operating with double floating point numbers and given the elapsed time, it comes out that the operations performed by our algorithm were $\simeq 15N < 6N$. This can be explained counting *all* of the operations performed during the cycle, included indexes operations and native counter increasing. Counting those operation, we reach $11N$, still below $15N$. The additional discrepancy can be explained taking into account memory access ($4N$ access operations) and interferences caused by the system processes.

5 Errors




Figure 3: Maximum percentage error of the numerical solution as a function of the number of points.

It is now interesting to analyse the error that we do in numerical approximation as a function of the number of points. It is intuitive that, as much as we restrict the step length h (that is the same thing as increasing the number of points) the error gets smaller and smaller. This behaviour is shown in Figure 3, where we plotted in a log-log scale the maximum percentage error calculated as

$$\epsilon_i = \left| \frac{v[i]}{u[i]} - 1 \right|$$

($v[i]$ is the numerical solution, $u[i]$ is the analytical one).

You see that the error lowers until $N = 10^5$, then increases again. This is a typical example of *loss of precision*; as our numerical solution gets closer to the analytical one, the ratio $v[i]/u[i]$ gets closer to one, and as a result in calculating ϵ_i we perform a subtraction between two almost identical (in our choice of precision) numbers. This causes a loss in terms of significant digits that explains the behaviour of the plot for big N . It is also evident that we get the lowest significant relative error for $N \simeq 10^6$. Relative errors for $N > 10^6$ are not worth trusting, due to the loss of precision explained above.

6 Conclusions

In this project we were able to develop an algorithm to solve a special differential equation using a numerical method but implementing it in various ways, which proved to be diversely capable. We were also able to observe a phenomenon of loss of precision in the computation of errors due to the use of the subtraction.

7 C++ Code

All the code is stored at https://github.com/matteosecli/Computational_Physics. Since this is a private repo, you have to request access writing to mattes@mail.uio.no.

```

#include <iostream>
#include <fstream>
#include <armadillo>
#include <cstdlib>
#include <cmath>
#include <ctime>
using namespace std;
using namespace arma;
namespace use {int onealg = 0; int out = 1;}
// 'fill_matrix' fills the matrix A in a tridiagonal form.
void fill_matrix(mat& A, int N) {
//Fill the matrix
A(0,0) = 2;
A(0,1) = -1;
for(int i = 1; i < N-1; i++){
A(i,i-1) = -1;
A(i,i) = 2;
A(i,i+1) = -1;
}
A(N-1,N-2) = -1;
A(N-1,N-1) = 2;
}
// 'solve_gaus' solves the system with a standard gaussian decomposition
void solve_gaus(vec& u, vec& f, int N){
//Start timing
clock_t t;
t = clock();
//Define our matrix and initialize it
mat A(N,N);
A.zeros();
fill_matrix(A,N);
//Just the decomposition
u = solve(A,f);
A.reset();
//Stop timing
t = clock() - t;
cout <<"Elapsed time (solve_gaus):\t\t" <<((float)t)/CLOCKS_PER_SEC <<"s." <<endl;
}
// 'solve_lu' solves the system with a standard LU decomposition
void solve_lu(vec& u, vec& f, int N){
//Start timing
clock_t t;
t = clock();
//Define our matrix and initialize it
mat A(N,N);
A.zeros();
fill_matrix(A,N);
//Define workspace matrices
mat L(N,N), U(N,N), P(N,N);
//Do the decomposition
lu(L, U, P, A);
A.reset();
//Just solve the system
vec b;
b = solve(L,P*f);
L.reset();
P.reset();
u = solve(U,b);
U.reset();
//Stop timing
t = clock() - t;
cout <<"Elapsed time (solve_lu):\t\t" <<((float)t)/CLOCKS_PER_SEC <<"s." <<endl;
}
// 'solvetrld' is a function that solves a linear sistem in 'u' relative to a
// tridiagonal matrix with diagonal elements equal to 'b', subdiagonal elements
// equal to 'a' and superdiagonal elements equal to 'c'. Returns the result
// in the vector 'u', overwriting its elements. The algorithm performs ~8N FLOPS.
void solvetrld(int& N, float& a, float& b, float& c, vec& u, vec& f){
// Start timing
clock_t t;

```

```

71 t = clock();
72 // Define the variable 'bet', that is just the denominator of 'gam',
73 // and 'gam' itself, that is a workspace vector
74 double bet;
75 vec gam(N);
76 // Start forward substitution
77 u[0]=f[0]/(bet=b);
78 for(int j = 1; j < N; j++) {
79     gam[j]=c/bet;
80     bet=b-a*gam[j];
81     u[j]=(f[j]-a*u[j-1])/bet;
82 }
83 // Just one-line backward substitution
84 for (int j = (N-2); j >= 0; j--) u[j] -= gam[j+1]*u[j+1];
85 // Stop timing and print elapsed time
86 t = clock() - t;
87 cout << "Elapsed time (solvetricid):\t\t" << ((float)t)/CLOCKS_PER_SEC << "s." << endl;
88 // Free space
89 gam.reset();
90 }
91 // 'solve_special' is a function that solves a special linear system
92 // relative to a tridiagonal matrix with b = 2 and a = c = -1. The
93 // solution has been found analytically, and once the pattern in
94 // the solution was recognized, it has been coded here. Warning! It
95 // overwrites 'u' and 'f', so make a copy before calling the function
96 // if you want to re-use them. The algorithm performs ~6N FLOPS.
97 void solve_special(int& N, vec& u, vec& f){
98     // Start timing
99     clock_t t;
100     t = clock();
101     for(int j = 1; j < N; j++) f[j]=(j+1)*f[j]+f[j-1];
102     u[N-1] = f[N-1]/(N+1);
103     int prev_idx = N;
104     for(int j = N - 1; j > 0; j--) {u[j-1]=(f[j-1]+j*u[j])/prev_idx; prev_idx = j;}
105     // Stop timing and print elapsed time
106     t = clock() - t;
107     cout << "Elapsed time (solve_special):\t\t" << ((float)t)/CLOCKS_PER_SEC << "s." << endl;
108 }
109 // 'split' is a function that discretizes the function 'func',
110 // storing its values in N points from 0 to 1 in the vector 'f'.
111 // Grid points are stored in the vector 'x'.
112 void split(vec& f, vec& x, int& N) {
113     // Define points spacing and calculate the grid points
114     // and the value of func in those points
115     double h = 1.0/(N+1);
116     double h_square = pow(h,2);
117     for(int i = 0; i < N; i++){
118         x[i] = (i+1)*h;
119         f[i] = h_square*100*exp(-10*x[i]);
120     }
121 }
122 // 'relative_error' calculates the relative error with respect to the
123 // theoretical value 'u_th(x)'.
124 vec relative_error(vec& u, vec& x, int& N) {
125     vec err(N);
126     err[0] = 0;
127     for(int i = 0; i <= N-1; i++){
128         err[i] = abs( u[i]/(1-(1-exp(-10))*x[i]-exp(-10*x[i])) - 1 );
129     }
130     return err;
131 }
132 // 'main' takes as first argumt the number of points that the program
133 // will use during the calculation. Use 'onealg 1' as second argument
134 // if you want to use only one algorithm and write to the output file.
135 // Use 'onealg 0' if you want to use only one algorithm and don't write
136 // to the output file.
137 int main(int argc, char *argv[])
138 {
139     int N = atoi(argv[1]);
140     // Perform some checks in the optional argument.
141     if(argc == 4) {
142         use::out = atoi(argv[3]);
143         if(strcmp(argv[2], "onealg") == 0 && (strcmp(argv[3], "1") == 0 || strcmp(argv[3], "0") == 0)) use::
144             onealg = 1;
145     }
146     else cout << "Wrong optional argument given. Use 'onealg 1' if you want to use only one algorithm (

```



```

    the fastest) and write to file; use 'onealg 0' if you instead want to write to the output file."
    << endl;
}
// Define the elements of the matrix related to the differential equation
float a = -1.0;
float b = 2.0;
float c = -1.0;
// Initialize the solution vector 'u' with zeros and the vector 'f'
// of the function values
vec u = zeros<vec>(N);
vec f(N), x(N);
//for(int i = 0; i < N; i++) f[i] = i+1;
// Discretize and define workspace vectors
split(f, x, N);
vec u_temp(N), f_temp(N), err(N);
// Compare algorithms only if we want to do benchmarks.
// This is to save memory if we want just to have grid numbers.
if(use::onealg == 0) {
    // Solve using 'solve_lu' and 'solve_gaus'
    if (N <= 10000) {
        solve_lu(u,f,N);
        solve_gaus(u,f,N);
    }
    // Solve using 'tridig'
    solvetrid(N, a, b, c, u, f);
}
// Solve using 'solve_special'
solve_special(N, u, f);
f.reset();
// Compute relative error only if 'onealg' is enabled, to speed up benchmarks
if(use::onealg == 1) {
    err = relative_error(u, x, N);
    cout << "Maximum relative error: " << err.max()*100 << "%" << endl;
}
// Write the resulting points on the output file
if(use::onealg == 1 && use::out == 1) {
    // Write the 'x' grid-points to the output file
    ofstream X;
    X.open("X.txt");
    X << x;
    X.close();
    // Write the 'u' grid-points to the output file
    ofstream U;
    U.open("U.txt");
    U << u;
    U.close();
    // Write the error bars to the output file
    ofstream E;
    E.open("E.txt");
    E << err;
    E.close();
}
return 0;
}

```